

Chapter 3: Process Concept

- The concept of Process
- Process Scheduling
- Operations on Processes
- Inter-process communication (IPC)


AE4B33OSS

Lecture 4 / Page 2

Silberschatz, Galvin and Gagne ©2005

Process Concept

- Textbooks use the terms *job* and *process* almost interchangeably
- Process – an instance of a program in execution;
 - Multiple instances of the same program are different processes
- A process is more than the program code(text section)
- A process also includes:
 - program counter
 - stack
 - data section
 - Heap
- Each process identified by a unique, positive integer id (process id)




AE4B33OSS

Lecture 4 / Page 3

Silberschatz, Galvin and Gagne ©2005

Process State

- As a process executes, it changes its **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a CPU
 - **terminated**: The process has finished execution



AE4B33OSS


Lecture 4 / Page 4

Silberschatz, Galvin and Gagne ©2005

Process Control Block (PCB)

Each process is represented in the OS by a PCB. Information associated with each process includes:

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information ("process environment")

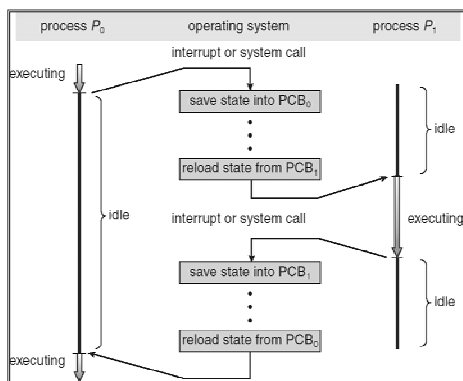


AE4B33OSS

Lecture 4 / Page 5

Silberschatz, Galvin and Gagne ©2005

CPU Switch From Process to Process



AE4B330SS

Lecture 4 / Page 6


Silberschatz, Galvin and Gagne ©2005

Process Representation in Linux

PCB in the Linux OS is Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this pro */
```

Doubly
linked
list



AE4B330SS

Lecture 4 / Page 7

Silberschatz, Galvin and Gagne ©2005

Threads

- A Thread, sometimes called lightweight process
- Thread are used for small tasks, whereas processes are used for more 'heavyweight' tasks – basically the execution of applications
- Both processes and threads are independent sequences of execution.
- Threads (of the same process) run in a shared memory space, while processes run in separate memory spaces.
- Single thread allows the process to perform one task at a time
- Modern OSs have extended the process concept to allow a process to have multiple threads of execution to perform more than one task at a time

AE4B330SS

Lecture 4 / Page 8

Silberschatz, Galvin and Gagne ©2005

Process Scheduling Queues

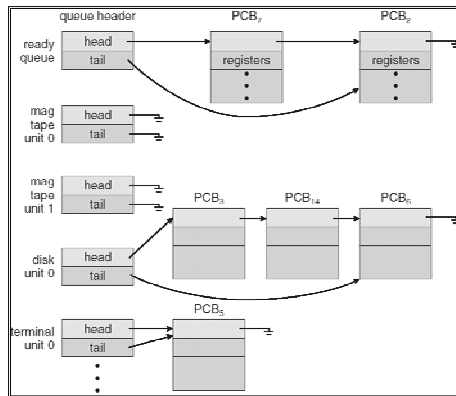
- **Job queue** – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues

AE4B330SS

Lecture 4 / Page 9

Silberschatz, Galvin and Gagne ©2005

Ready Queue and Various I/O Device Queues

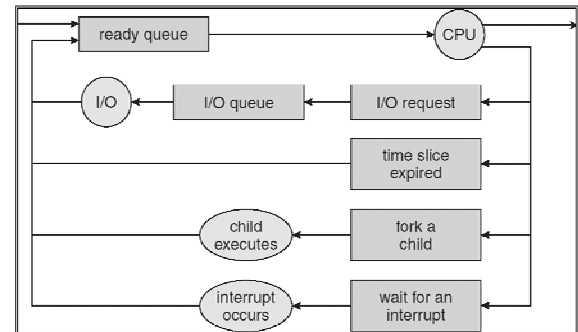


AE4B330SS

Lecture 4 / Page 10

Silberschatz, Galvin and Gagne ©2005

Simplified Model of Process Scheduling



AE4B330SS

Lecture 4 / Page 11

Silberschatz, Galvin and Gagne ©2005

Schedulers

■ Long-term scheduler (or job scheduler) – selects which processes should be brought into the ready queue

- Long-term scheduler is invoked very infrequently (seconds, minutes) ⇒ (may be slow)
- The long-term scheduler controls the *degree of multiprogramming*

■ Short-term scheduler (or CPU scheduler) – selects which process should be executed next and allocates CPU for that process

- Short-term scheduler is invoked very frequently (milliseconds) ⇒ (must be fast)

■ Processes can be described as either:

- **I/O-bound process** – spends more time doing I/O than computations
- **CPU-bound process** – spends more time doing computations


AE4B330SS

Lecture 4 / Page 12

Silberschatz, Galvin and Gagne ©2005

Schedulers

- **Medium-term scheduler** : To remove processes from memory and later, the process can be reintroduced into memory and its execution can be continued where it left off. This scheme is called swapping



Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**.
- **Context** of a process represented in the PCB
- Context-switch time is *overhead*; the system does no useful work while switching
- Switch time varies from system to system – dependent on hardware support(memory speed, registers)

AE4B330SS

Lecture 4 / Page 14

Silberschatz, Galvin and Gagne ©2005

Process Creation

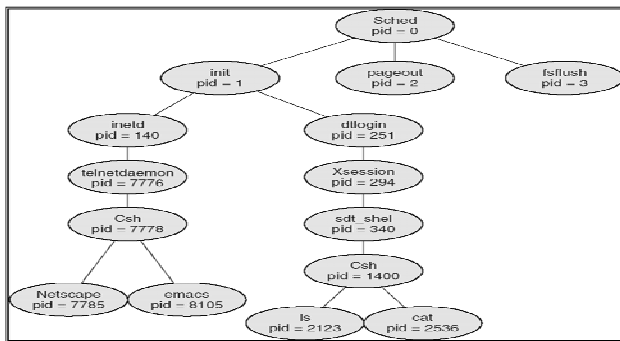
- Parent process create children processes
 - Children, in turn create other processes, forming a tree of processes
- Resource sharing possibilities
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution possibilities
 - Parent and children can execute concurrently, or
 - Parent can wait until some or all its children terminate
- Memory address space possibilities
 - Address space of child duplicate of parent
 - Child has a new program loaded into it
- POSIX examples
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program

AE4B330SS

Lecture 4 / Page 15

Silberschatz, Galvin and Gagne ©2005

Process Creation Illustrated



A Tree of processes on a typical Solaris System

AE4B330SS

Lecture 4 / Page 16

Silberschatz, Galvin and Gagne ©2005

Process Creation

- A process can create another process by invoking **fork()** system call
 - The process that invokes the **fork()** is known as the **parent** and the new process is called the **child**
- NOTE: The **fork** system call does not take an argument.

AE4B330SS

Lecture 4 / Page 17

Silberschatz, Galvin and Gagne ©2005

A Simple program of fork()

```
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main(void)
{
    printf("Hello \n");
    fork();
    printf("bye\n");
    return 0;
}
```

A summary of fork() return values follows:
 fork_return > 0: this is the parent
 fork_return == 0: this is the child
 fork_return == -1: fork() failed and there is no child

AE4B330SS

Lecture 4 / Page 18

Silberschatz, Galvin and Gagne ©2005

Creating a Separate Process Using fork() system call


```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        printf("Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

AE4B330SS

Lecture 4 / Page 19

Silberschatz, Galvin and Gagne ©2005

Process creation using fork() system call



AE4B330SS

Lecture 4 / Page 20

Silberschatz, Galvin and Gagne ©2005

Process Termination

- A process terminates when it finishes executing its last statement and asks the operating system to terminate it using exit() system call
- Process encounters a fatal error
 - Can be for many reasons like arithmetic exception etc.
- Parent may terminate execution of children processes using kill() system call. *Some possible reasons*
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
- Parent is exiting
 - Some operating systems may not allow child to continue if its parent terminates

AE4B330SS

Lecture 4 / Page 21

Silberschatz, Galvin and Gagne ©2005

Cooperating Processes

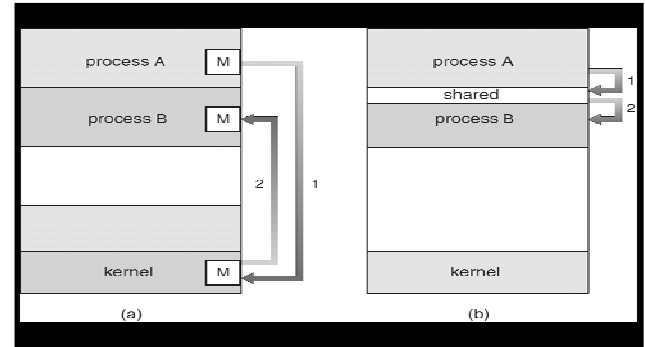
- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience
- **Producer-Consumer Problem**
 - Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size

AE4B330SS

Lecture 4 / Page 22

Silberschatz, Galvin and Gagne ©2005

Interprocess Communication (IPC)



Two models of IPC: (a) Message passing (b) Shared memory

AE4B330SS

Lecture 4 / Page 23

Silberschatz, Galvin and Gagne ©2005

IPC

Message Passing	Shared Memory
Easier to implement, particularly useful in distributed environment	Not easy to implement
Not faster	Faster
Implemented using series of system calls and more time consuming	Systems calls are required only to establish shared-memory regions
No concurrent execution	Processes runs concurrently


AE4B330SS

Lecture 4 / Page 24

Silberschatz, Galvin and Gagne ©2005

Shared Memory IPC

- Processes share a common piece of memory either physically or virtually
- Communications via normal reads/writes



AE4B330SS

Lecture 4 / Page 25

Silberschatz, Galvin and Gagne ©2005

Shared memory

▪Eg: Producer Consumer Problem

- Producer produces items, consumer consumes item
 - Unbounded Buffer
 - Bounded Buffer
- Shared buffer implemented as a circular array with two logical pointers in and out

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0; //next free position
int out = 0; //first full position
```

AE4B330SS

Lecture 4 / Page 26

Silberschatz, Galvin and Gagne ©2005

Code for Consumer Process

```
item nextProduced;
while (true) {
    /* Produce an item in nextProduced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}
```

The producer process

AE4B330SS

Lecture 4 / Page 27

Silberschatz, Galvin and Gagne ©2005

Code Consumer Process

```
item nextConsumed;
while (true) {
    while (in == out)
        ; // do nothing -- nothing to consume

    // remove an item from the buffer
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in nextConsumed */
}
```

The consumer process

AE4B330SS

Lecture 4 / Page 28

Silberschatz, Galvin and Gagne ©2005

Message-Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- A message passing system provides two operations:
 - **send**(message) – message size fixed or variable
 - **receive**(message)
- If *P* and *Q* wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., logical properties)

AE4B330SS

Lecture 4 / Page 29

Silberschatz, Galvin and Gagne ©2005

Message Passing

■ Methods for implementing the communication link and the *send()*/*receive()* operations:

1. Direct or Indirect communication (*Naming*)
2. Synchronous or Asynchronous communication
3. Automatic or Explicit buffering

Process A	Process B
<pre>while (TRUE) { produce an item send (B, item) }</pre>	<pre>while (TRUE) { receive (A, item) consume item }</pre>

AE4B33OSS

Lecture 4 / Page 30

Silberschatz, Galvin and Gagne ©2005

Direct Communication

■ Direct Communication

- Processes must name each other explicitly:
 - **Symmetric Addressing**
 - *send (P, message)* – send a message to process P
 - *receive(Q, message)* – receive a message from process Q
 - **Asymmetric addressing**
 - *send (P, message)* – send to process P
 - *receive(id, message)* – rx from any; system sets id = sender
- Properties of communication link
 - Links are established automatically between pairs
 - A link is associated with exactly two processes
 - Between each pair of processes, there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional
- **Disadvantage:** a process must know the name or ID of the process(es) it wishes to communicate with

AE4B33OSS

Lecture 4 / Page 31

Silberschatz, Galvin and Gagne ©2005

Indirect Communication

■ Indirect Communication

- Messages are directed and received from *mailboxes* (also referred to as *ports*)
 - Each mailbox has a unique *id* and is created by the kernel on request
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

AE4B33OSS


Lecture 4 / Page 32

Silberschatz, Galvin and Gagne ©2005

Indirect Communication

■ Operations

- create a new mailbox
- send and receive messages mailbox
- destroy a mailbox



■ Primitives are defined as:

- *send(A, message)* – send a message to mailbox A.
- *receive(A, message)* – receive a message from mailbox A.

A. Frank - P. Weisberg
Lecture 4 / Page 33

Silberschatz, Galvin and Gagne ©2005

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A.
 - P_1 sends; P_2 and P_3 receive.
 - Who gets the message?
- Possible solutions:
 - Allow a link to be associated with at most two processes.
 - Allow only one process at a time to execute a receive operation.
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



A. Frank - P. Weisberg
Lecture 4 / Page 34

Silberschatz, Galvin and Gagne ©2005

Synchronization

- Message passing may be either blocking or non-blocking also known as synchronous and asynchronous
- **Blocking** is considered **synchronous**
 - **Blocking send**: the sender blocks until the message is received by the other party
 - **Blocking receive**: the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send**: the sender sends the message and continues executing
 - **Non-blocking receive**: the receiver gets either a valid message or a null message (when nothing has been sent to the receiver)
- Often a combination:
 - Non-blocking send and blocking receive

AE4B330SS

Lecture 4 / Page 35

Silberschatz, Galvin and Gagne ©2005

Link Capacity – Buffering

- Queue of messages attached to the link; implemented in one of three ways:
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous).
 2. Bounded capacity – finite length of n messages
Sender must wait if link full.
 3. Unbounded capacity – infinite length
Sender never waits.

AE4B330SS

Lecture 4 / Page 36

Silberschatz, Galvin and Gagne ©2005