

Multithreaded Programming

- Overview
- Multithreading Models
- Thread Libraries
- Threading Issues

AE4B330SS

Lecture 4 / Page 37

Silberschatz, Galvin and Gagne ©2005

What are Threads?


- Thread
 - Independent stream of instructions
 - Basic unit of CPU utilization
- A thread contains
 - A thread ID
 - A register set (including the Program Counter PC)
 - An execution stack
- A thread shares with its sibling threads
 - The code, data and heap section
 - Other OS resources, such as open files and signals

AE4B330SS

Lecture 4 / Page 38

Silberschatz, Galvin and Gagne ©2005

Process Address Space Revisited



(a) Process with Single Thread


(b) Process with Two Threads

AE4B330SS

Lecture 4 / Page 39

Silberschatz, Galvin and Gagne ©2005

Single and Multithreaded Processes



single-threaded process

multithreaded process

AE4B330SS

Lecture 4 / Page 40

Silberschatz, Galvin and Gagne ©2005

Multi-Threaded Processes

- Each thread has a private stack
- But threads share the process address space!
- There's no memory protection!
- Threads could potentially write into each other's stack

AE4B330SS

Lecture 4 / Page 41

Silberschatz, Galvin and Gagne ©2005

Why use Threads?

- A specific example, a Web server:

```
do
{
    get web page request from client
    check if page exists and client has permissions
    transmit web page back to client
} while(1);
```

- If transmission takes very long time, server is unable to answer other client's requests. Solution:


```
do
{
    get web page request from client
    check if page exists and client has permissions
    create a thread to transmit web page back to client
} while(1);
```

AE4B330SS

Lecture 4 / Page 42

Silberschatz, Galvin and Gagne ©2005

Multithreaded Server Architecture



AE4B330SS

Lecture 4 / Page 43

Silberschatz, Galvin and Gagne ©2005

Benefits of Multithreaded Programming

- Responsiveness
- Resource Sharing
- Economy
- Scalability

AE4B330SS

Lecture 4 / Page 44

Silberschatz, Galvin and Gagne ©2005

Multicore Programming

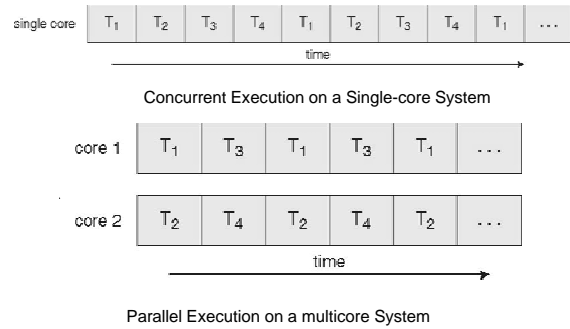
- Multicore systems putting pressure on programmers, challenges include
 - Dividing activities
 - Balance
 - Data splitting
 - Data dependency
 - Testing and debugging

AE4B33OSS

Lecture 4 / Page 45

Silberschatz, Galvin and Gagne ©2005

Multicore Programming



AE4B33OSS

Lecture 4 / Page 46

Silberschatz, Galvin and Gagne ©2005

User Threads

- Thread management done by user-level threads library
- Three primary thread libraries:
 - POSIX Pthreads
 - Win32 threads
 - Java threads

AE4B33OSS

Lecture 4 / Page 47

Silberschatz, Galvin and Gagne ©2005

Kernel Threads

- Supported by the Kernel
- Examples
 - Windows XP/2000
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X

AE4B33OSS

Lecture 4 / Page 48

Silberschatz, Galvin and Gagne ©2005

Multithreading Models

- Relationships between user threads and kernel threads
 - Many-to-One
 - One-to-One
 - Many-to-Many

AE4B330SS

Lecture 4 / Page 49

Silberschatz, Galvin and Gagne ©2005

Many-to-One


- Many user-level threads mapped to single kernel thread
 - Thread management by thread library in user space → efficient
 - No concurrency
 - Unable to run in parallel on multiprocessors (MP)
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads

AE4B330SS

Lecture 4 / Page 50

Silberschatz, Galvin and Gagne ©2005

Many-to-One Model



AE4B330SS

Lecture 4 / Page 51

Silberschatz, Galvin and Gagne ©2005

One-to-One


- Each user-level thread maps to a kernel thread
 - More concurrency
 - Multiple threads to run in parallel on multiprocessors
 - Overhead of creating kernel threads
- Examples
 - Windows NT/XP/2000
 - Linux
 - Solaris 9 and later

AE4B330SS

Lecture 4 / Page 52

Silberschatz, Galvin and Gagne ©2005

One-to-one Model



AE4B330SS

Lecture 4 / Page 53

Silberschatz, Galvin and Gagne ©2005

Many-to-Many Model


- Allows many user level threads to be mapped to a smaller or equal number of kernel threads
 - Concurrency
 - Can run in parallel on a multiprocessor
 - Developers can create as many user threads as necessary
- Examples
 - Solaris prior to version 9
 - Windows NT/2000 with the *ThreadFiber* package

AE4B330SS

Lecture 4 / Page 54

Silberschatz, Galvin and Gagne ©2005

Many-to-Many Model



AE4B330SS

Lecture 4 / Page 55

Silberschatz, Galvin and Gagne ©2005

Two-level Model


- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier

AE4B330SS

Lecture 4 / Page 56

Silberschatz, Galvin and Gagne ©2005

Two-level Model



AE4B330SS

Lecture 4 / Page 57

Silberschatz, Galvin and Gagne ©2005

Thread Libraries

- Thread libraries provide the programmer with an API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

AE4B330SS

Lecture 4 / Page 58

Silberschatz, Galvin and Gagne ©2005

Pthreads

- Refers to the POSIX standard (IEEE 1003.1c)
- API for thread creation and synchronization
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)
- May be provided either as user-level or kernel-level

AE4B330SS

Lecture 4 / Page 59

Silberschatz, Galvin and Gagne ©2005

Java Threads

- Java threads may be created by:
 - Extending Thread class
 - Implementing the Runnable interface
- JVM manages Java threads
 - Creation
 - Execution
 - Etc.

AE4B330SS

Lecture 4 / Page 60

Silberschatz, Galvin and Gagne ©2005

Practice Exercise

- Consider two threads sharing a global variable count, initially 10:

Thread A <div style="border: 1px solid black; padding: 2px; display: inline-block;">count++;</div>	Thread B <div style="border: 1px solid black; padding: 2px; display: inline-block;">count--;</div>
--	--

- What are the possible values for count after both threads finish executing:

AE4B330SS

Lecture 4 / Page 61

Silberschatz, Galvin and Gagne ©2005

Practice Exercise

- Consider the following three concurrent threads that share a global variable g, initially 10:

Thread A <div style="border: 1px solid black; padding: 2px; display: inline-block;">g = g * 2;</div>	Thread B <div style="border: 1px solid black; padding: 2px; display: inline-block;">g = g + 1;</div>	Thread C <div style="border: 1px solid black; padding: 2px; display: inline-block;">g = g - 2;</div>
--	--	--

- What are the possible values for g, after all three threads finish executing?

AE4B330SS

Lecture 4 / Page 62

Silberschatz, Galvin and Gagne ©2005

Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Thread cancellation
- Signal handling
- Thread pools
- Thread specific data
- Scheduler activations

AE4B330SS

Lecture 4 / Page 63

Silberschatz, Galvin and Gagne ©2005

Semantics of fork() and exec()

- Does **fork()** duplicate only the calling thread or all threads?
 - Which of the two versions of fork() to use depends on the application
 - exec() immediately after fork()
 - No exec() after fork()

AE4B330SS

Lecture 4 / Page 64

Silberschatz, Galvin and Gagne ©2005

Thread Cancellation

- Terminating a thread before it has completed
- A thread that is cancelled is called target thread
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- Problems: when resources have been allocated to a canceled thread or while in the midst of updating data sharing with other threads

AE4B33OSS

Lecture 4 / Page 65

Silberschatz, Galvin and Gagne ©2005

Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred
 - A **signal handler** is used to process signals. All signals follow this pattern
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Once delivered, the Signal must be handled
- ▶ Synchronous vs. asynchronous signals
- ▶ A signal may be handled by one of following handlers:
 - ▶ Default signal handlers
 - ▶ user-defined handlers

AE4B33OSS

Lecture 4 / Page 66

Silberschatz, Galvin and Gagne ©2005

Signal Handling

- Options to deliver signals in multithreaded programs:
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process
- ex:
 - UNIX function for delivering a signal is `kill(pid_t pid, int signal)`
 - POSIX thread provides the `pthread_kill(pthread_t tid, int signal)`

AE4B33OSS

Lecture 4 / Page 67

Silberschatz, Galvin and Gagne ©2005

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
- Ex:
 - `QueueUserWorkItem(LPTHREAD_START_ROUTINE Function, PVOID Param, ULONG Flags)` in Win32 API
 - `java.util.concurrent` package in Java 1.5

AE4B33OSS

Lecture 4 / Page 68

Silberschatz, Galvin and Gagne ©2005

Thread-Specific Data

- Allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)

AE4B330SS

Lecture 4 / Page 69

Silberschatz, Galvin and Gagne ©2005

Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the thread library
 - An intermediate data structure called *LWP* (Lightweight Process) between user thread and kernel thread
- This communication allows an application to maintain the correct number of kernel threads

AE4B330SS

Lecture 4 / Page 70

Silberschatz, Galvin and Gagne ©2005

Operating System Examples

- Windows XP Threads
- Linux Threads

AE4B330SS

Lecture 4 / Page 71

Silberschatz, Galvin and Gagne ©2005

Windows XP Threads

- Implements the one-to-one mapping
- Each thread contains
 - A thread id
 - Register set
 - Separate user and kernel stacks
 - Private data storage area
- The primary data structures of a thread include:
 - ETHREAD (executive thread block)
 - KTHREAD (kernel thread block)
 - TEB (thread environment block)

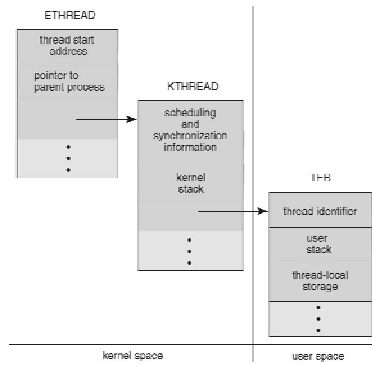
Context of the threads

AE4B330SS

Lecture 4 / Page 72

Silberschatz, Galvin and Gagne ©2005

Windows XP Threads



AE4B330SS

Lecture 4 / Page 73

Silberschatz, Galvin and Gagne ©2005

Linux Threads

- Linux refers to them as *tasks* rather than process or *thread*
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process)
 - CLONE_FS, CLONE_VM, CLONE_SIGHAND, CLONE_FILES

AE4B330SS

Lecture 4 / Page 74

Silberschatz, Galvin and Gagne ©2005

Linux Threads

- Linux provides clone() system call to create threads
 - When clone() is invoked, it is passed set of flags, which determine how much sharing is to take place between parent and child tasks


flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

AE4B330SS

Lecture 4 / Page 75

Silberschatz, Galvin and Gagne ©2005

Java Thread States



AE4B330SS

Lecture 4 / Page 76

Silberschatz, Galvin and Gagne ©2005