

## Inter-process Communication and Synchronization



## Contents

- Background
- The Critical-Section Problem
- Peterson's Solutions
- Synchronization Hardware
- Semaphores
- Classical Synchronization Tasks
- Monitors

## Background

- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution
- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

## Producer-Consumer Problem (PCP)

- Paradigm for cooperating processes
- *producer process produces information that is consumed by a consumer process.*
  - *unbounded-buffer places no practical limit on the size of the buffer.*
  - *bounded-buffer assumes that there is a fixed buffer size.*
- Basic synchronization requirement
  - Producer should not write into a full buffer
  - Consumer should not read from an empty buffer
  - All data written by the producer must be read exactly once by the consumer

## Bounded-Buffer – Shared-Memory Solution to PCP

- The shared buffer is implemented as a circular array with two pointers: in and out (in: points to next free position in the buffer; out: points to first full position in the buffer)

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

## Bounded-Buffer: Producer Process

```
item nextProduced;
while (1) {
    /* Produce an item in nextProduced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
}
```

**The producer process**

## Bounded-Buffer: Consumer Process

```
item nextConsumed;
while (1) {
    while (in == out)
        ; // do nothing -- nothing to consume

    // remove an item from the buffer
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in nextConsumed */
}
```

**The consumer process**

## Bounded-Buffer – Shared-Memory Solution to PCP

- This solution allows at most BUFFER\_SIZE-1 items in the buffer (of size BUFFER\_SIZE) at the same time. A solution, where BUFFER\_SIZE items can be in the buffer is *not simple*
- Suppose we modify the producer-consumer code by adding a variable *counter*, initialized to 0 and incremented each time a new item is added to the buffer

## Bounded-Buffer for BUFFER\_SIZE items

### ■ Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;
item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int counter = 0;
```

## Bounded-Buffer for BUFFER\_SIZE items

### ■ Producer process

```
item nextProduced;

while (1) {
    while (counter == BUFFER_SIZE)
        ; /* do nothing */
    buffer[in] = nextProduced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

### ■ Consumer process

```
item nextConsumed;

while (1) {
    while (counter == 0)
        ; /* do nothing */
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
}
```

## The Problem with this solution

- The statement “**counter++**” may be implemented in machine language as:

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- The statement “**counter--**” may be implemented as:

```
register2 = counter
register2 = register2 - 1
counter = register2
```

## The Problem with this solution

- If both the producer and consumer attempt to update *counter concurrently*, the assembly language statements may get interleaved.
- Interleaving depends upon how the producer and consumer processes are scheduled

### The Problem with this solution

- Assume **counter** is initially 5. One interleaving of statements is:

$T_0$ : producer: **register1 = counter** (*register1 = 5*)  
 $T_1$ : producer: **register1 = register1 + 1** (*register1 = 6*)  
 $T_2$ : consumer: **register2 = counter** (*register2 = 5*)  
 $T_3$ : consumer: **register2 = register2 - 1** (*register2 = 4*)  
 $T_4$ : producer: **counter = register1** (*counter = 6*)  
 $T_5$ : consumer: **counter = register2** (*counter = 4*)

- The value of **count** may be either 4 or 6, where the correct result should be 5.

AE4B330SS

Lecture 6 / Page 13

Silberschatz, Galvin and Gagne ©2005

### Race Condition

- A scenario in which several processes may read and write shared data concurrently and the final value of the shared data depends upon which process finishes last
  - So the final output is dependent on the relative speed of the processes
- Race conditions must be prevented
  - concurrent processes must be synchronized

AE4B330SS

Lecture 6 / Page 14

Silberschatz, Galvin and Gagne ©2005

### Atomic Operation

- An operation that is either executed fully without interruption, or not executed at all
  - The "operation" can be a group of instructions
  - Ex. the instructions for *counter++* and *counter--*
- Note that the producer-consumer problem's solution works if *counter++* and *counter--* are made atomic
- In practice, the process may be interrupted in the middle of an atomic operation, but the atomicity should ensure that no process uses the effect of the partially executed operation until it is completed

AE4B330SS

Lecture 6 / Page 15

Silberschatz, Galvin and Gagne ©2005

### The Critical-Section Problem

- $n$  processes all competing to use some shared data
- Each process has a code segment, called *critical section*, in which the shared data is accessed.
- Critical-Section Problem – ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.

AE4B330SS

Lecture 6 / Page 16

Silberschatz, Galvin and Gagne ©2005

### Requirements for Solution to Critical-Section Problem

1. **Mutual Exclusion.** If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections.
2. **Progress.** If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely.
3. **Bounded Waiting/No Starvation.** A bound (or limit) must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.

AE4B330SS

Lecture 6 / Page 17

Silberschatz, Galvin and Gagne ©2005

### Critical Section

- Entry section: a piece of code executed by a process just before entering a critical section
- Exit section: a piece of code executed by a process just after leaving a critical section
- The remaining code is the *remainder section*
- General structure of process  $P_i$  is

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

AE4B330SS

Lecture 6 / Page 18

Silberschatz, Galvin and Gagne ©2005

### Peterson's Solution

- The Peterson's solution restricted to two processes  $P_0$  and  $P_1$  that alternate execution between their critical sections and remainder sections.
- The two processes  $P_i$  and  $P_j$  ( $j = 1-i$ ) share two variables:
  - `int turn == 0;`
  - `boolean flag[2]`
- The variable `turn` indicates whose turn it is to enter the critical section. The variable `turn` is set to either 0 or 1 randomly (or it can always be set to 0)
  - `turn == i` implies  $P_i$ 's turn to enter its critical section
- The flag array is used to indicate if a process is ready to enter the critical section.
  - initially, `flag[0] == flag[1] == false`
  - `flag[i] == true` implies that  $P_i$  is ready to enter its critical section

AE4B330SS

Lecture 6 / Page 19

Silberschatz, Galvin and Gagne ©2005

### Algorithm for Process $P_i$ in Peterson's solution

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while ( flag[j] && turn == j);  
    critical section  
    flag[i] = FALSE;  
    remainder section  
} while(TRUE);
```

- Provetthat that
  1. Mutual exclusion is preserved
  2. Progress requirement is satisfied
  3. Bounded-waiting requirement is met

AE4B330SS

Lecture 6 / Page 20

Silberschatz, Galvin and Gagne ©2005

## Synchronization Hardware

- Many systems provide hardware support for critical section code
- A process must acquire lock before entering a critical section; it releases the lock when it exits the critical section

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

AE4B330SS

Lecture 6 / Page 21

Silberschatz, Galvin and Gagne ©2005

## Synchronization Hardware

- The critical-section problem could be solved in a uniprocessor environment if we could prevent interrupts from occurring while shared variable was being modified
- Unfortunately, this solution is not as feasible in a multiprocessor environment
  - Disabling interrupts on a multiprocessor can be time consuming, as the message passed to all the processors
  - Message passing delays entry into each critical section, & system efficiency decreases
  - The effect on system's clock if the clock is kept updated by interrupts

AE4B330SS

Lecture 6 / Page 22

Silberschatz, Galvin and Gagne ©2005

## Hardware Instruction Based Solutions

- TestAndSet: Test and modify the content of a word atomically
- The TestAndSet() instruction is defined below:

```
boolean TestAndSet (boolean *target)  
{  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

AE4B330SS

Lecture 6 / Page 23

Silberschatz, Galvin and Gagne ©2005

## Mutual-exclusion with TestAndSet()

- Shared data:  
boolean lock = false
- Process  $P_i$

```
do {  
    while(TestAndSet(&lock));  
    critical section  
    lock = FALSE;  
    remainder section  
} while(TRUE);
```

AE4B330SS

Lecture 6 / Page 24

Silberschatz, Galvin and Gagne ©2005

### Swap() instruction

- Swap: Atomically swap content of two words.
- The Swap() instruction is defined as follows:

```
void Swap(boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

AE4B330SS

Lecture 6 / Page 25

Silberschatz, Galvin and Gagne ©2005

### Mutual-exclusion with swap() instruction

- Uses two Boolean variables lock and Key
- If lock == false then a process can enter the critical section, otherwise it can't
- Initially lock == false
- Process P<sub>i</sub>

```
do {
    key = TRUE;
    while ( key == TRUE)
        Swap (&lock, &key );
    critical section
    lock = FALSE;
    remainder section
} while (TRUE);
```

AE4B330SS

Lecture 6 / Page 26

Silberschatz, Galvin and Gagne ©2005

### Bounded-waiting mutual exclusion with TestAndSet()

- Shared data (initialized to false)
  - boolean waiting[n];
  - boolean lock;
- We present another algorithm using TestAndSet() that satisfies all the critical-section requirements

AE4B330SS

Lecture 6 / Page 27

Silberschatz, Galvin and Gagne ©2005

### Bounded-waiting Mutual Exclusion with TestAndSet ()

```
do {
    waiting[i] = TRUE;
    key = TRUE;
    while (waiting[i] && key)
        key = TestAndSet(&lock); //only first lock==false will set key=false
    waiting[i] = FALSE;
    critical section
    j = (i + 1) % n; //look for the next P[j] waiting: bound- waiting req.
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = FALSE; //wakeup only one process P[j] without releasing lock
    remainder section
} while (TRUE);
```

AE4B330SS

Lecture 6 / Page 28

Silberschatz, Galvin and Gagne ©2005

### Bounded-waiting mutual exclusion with TestAndSet()

- Mutual exclusion
  - Mutual exclusion requirement is met if either waiting[i] == false or key == false
- Progress
  - Exit process sends a process in
- Bounded waiting
  - Any process wait at most n-1 times to enter critical section
- Atomic TestAndSet is hard to implement in multiprocessor environment

AE4B330SS

Lecture 6 / Page 29

Silberschatz, Galvin and Gagne ©2005

### Semaphores

- The hardware-based solutions to the critical-section problem are complicated for application programmers to use. To overcome this difficulty, we can use a synchronization tool called a semaphore
- Semaphores
  - A variable used for signaling between the process
  - Originally semaphores were flags for signaling between ships
- Semaphore  $S$  – integer variable
- Two standard operations modify  $S$ :
  - wait() (or acquire)
  - signal() (or release)
- These two operations are indivisible (**atomic**)

AE4B330SS

Lecture 6 / Page 30

Silberschatz, Galvin and Gagne ©2005

### Semaphore Operations

- The definition of wait () is as follows

```
wait (S) {  
    while (S <= 0)  
        ; // no-op, called busy waiting, spinlock  
    S--;  
}
```

- The definition of signal() is as follows:

```
signal (S) {  
    S++;  
}
```

- Busy waiting: While a process is in its critical section, any other process that tries to enter its critical section must loop continuously in the entry code
- Spinlock: the process “spins” while waiting for the lock
- Must guarantee that no two processes can execute wait () and signal () on the same semaphore at the same time

AE4B330SS

Lecture 6 / Page 31

Silberschatz, Galvin and Gagne ©2005

### Semaphore

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1; can be simpler to implement
  - Also known as **mutex locks (are the locks that provide mutual exclusion)**
- Can implement a counting semaphore  $S$  as a binary semaphore

AE4B330SS

Lecture 6 / Page 32

Silberschatz, Galvin and Gagne ©2005



### Mutual-exclusion with semaphores for n Processes

- Let the n processes share a semaphore, mutex, initialized to 1
- Process  $P_i$

```
do {
    wait (mutex);
    // Critical Section
    signal (mutex);
    // remainder section
} while (TRUE);
```

AE4B330SS

Lecture 6 / Page 33

Silberschatz, Galvin and Gagne ©2005

### Semaphore Usage

- Consider two concurrently running processes:  $P_1$  with a statement  $S_1$  and  $P_2$  with a statement  $S_2$ .
- Suppose we require that  $S_2$  be executed only after  $S_1$  has completed
- Let  $P_1$  and  $P_2$  share a common semaphore synch, initialized to zero.

```
P1:
    S1;
    signal (synch);
P2:
    wait (synch);
    S2;
```

- **Process synchronization:**  $S_2$  in  $P_2$  should be executed after  $S_1$  in  $P_1$

AE4B330SS

Lecture 6 / Page 34

Silberschatz, Galvin and Gagne ©2005

### Semaphore Implementation with no Busy waiting

- Define semaphore as a "C" structure

```
typedef struct {
    int value;
    struct process *list;
} semaphore;
```

- Two operations:
  - The **block ()** operation suspends the process that invokes it.
  - The **wakeup(P)** operation resumes the execution of a blocked process P.

AE4B330SS

Lecture 6 / Page 35

Silberschatz, Galvin and Gagne ©2005

### Semaphore Implementation with no Busy waiting

- The wait() operation is now defined as

```
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}
```

- The signal() operation is now defined as

```
signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

AE4B330SS

Lecture 6 / Page 36

Silberschatz, Galvin and Gagne ©2005

## Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
  - example: process  $P_0$  waits for process  $P_1$ , and process  $P_1$  waits for process  $P_0$
- Consider two processes,  $P_0$  and  $P_1$ , each accessing 2 semaphores,  $S$  and  $Q$  initialized to 1

$P_0$	$P_1$	
wait(S); //exec 1st	wait(Q); //exec 2nd	$P_0$ is waiting for $P_1$ to execute signal (Q)
wait(Q); //exec 3rd	wait(S); //exec 4th	$P_1$ is waiting for $P_0$ to execute signal (S)
...	...	
signal(S);	signal(Q);	
signal(Q);	signal(S);	Both Processes $P_0$ and $P_1$ are in a deadlock

- **Starvation or indefinite blocking**
  - A situation in which processes wait indefinitely within the semaphore. It may occur if we remove processes in LIFO order

AE4B330SS

Lecture 6 / Page 37

Silberschatz, Galvin and Gagne ©2005

## Priority Inversion Problem

- Scheduling challenge arises when lower-priority process holds a lock needed by higher-priority process
- Scenario :
  - Processes with priorities  $L < M < H$  request resource  $R$
  - $L$  first locks on  $R$ ; then  $H$  requests  $R$ , but normally  $H$  would wait for  $L$  to finish using resource  $R$ .
  - In the meantime,  $M$  requests  $R$  and preempts  $L$
  - $H$  is waiting longer for lower-priority processes
- Not a problem with only two levels of priorities, but two levels are insufficient for most OSs
- Solved via **priority-inheritance protocol**
  - All lower-priority processes that are accessing resources requested by a higher-priority process, inherits the higher-priority level until it releases the resource

AE4B330SS

Lecture 6 / Page 38

Silberschatz, Galvin and Gagne ©2005

## Classical Problems of Synchronization

- Bounded-Buffer Producer-Consumer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

AE4B330SS

Lecture 6 / Page 39

Silberschatz, Galvin and Gagne ©2005

## Bounded-Buffer Producer-Consumer Problem


- Two processes (one producer, one consumer) share a common, fixed-size buffer
- Producer places information into the buffer
- Consumer takes it out
- The producer wants to place a new item into the buffer, but the buffer is already full
  - Solution: producer goes to sleep, wakes up when the consumer has emptied one or more items
- Consumer wants to consume an item, but the buffer is empty
  - Solution: consumer goes to sleep, wakes up when the producer has produced items

AE4B330SS

Lecture 6 / Page 40

Silberschatz, Galvin and Gagne ©2005

## Producer Consumer Problem



6

AE4B33OSS

Lecture 6 / Page 41

Silberschatz, Galvin and Gagne ©2005

## Bounded-Buffer Producer-Consumer Problem

- The binary semaphore mutex provides mutually exclusive access to the buffer & is initialized to the value 1
- The counting semaphore full used for counting the number of slots that are full & is initialized to 0
- The counting semaphore empty used for counting the number of slots that are empty & is initialized to n, the buffer size
- The empty and full semaphores are used for process synchronization

AE4B33OSS

Lecture 6 / Page 42

Silberschatz, Galvin and Gagne ©2005

## Bounded-Buffer Producer-Consumer Problem

- The structure of the producer process

```
do {
    ...
    // produce an item
    wait (empty); //wait for an empty space
    wait (mutex);
    ...
    // add the item to the buffer
    ...
    signal (mutex);
    signal (full); // report the new full slot
} while(TRUE);
```

AE4B33OSS

Lecture 6 / Page 43

Silberschatz, Galvin and Gagne ©2005

## Bounded-Buffer Producer-Consumer Problem

- The structure of the consumer process

```
do {
    wait (full); // wait for a stored item
    wait (mutex);
    ...
    // remove an item from buffer
    ...
    signal (mutex);
    signal (empty); //report the new empty slot
    ...
    // consume the removed item
    ...
} while (TRUE)
```

AE4B33OSS

Lecture 6 / Page 44

Silberschatz, Galvin and Gagne ©2005

## Bounded-Buffer Producer-Consumer Problem

<b>Producer</b>	<b>Consumer</b>
produce an item wait(empty) <ul style="list-style-type: none"><li>• loose processor</li></ul>	
wait(mutex) add item to buffer signal(mutex) signal(full)	wait(full) % cannot proceed because there are no items in the buffer <ul style="list-style-type: none"><li>• loose processor</li></ul>
produce an item wait(empty) wait(mutex) add item to buffer signal(mutex) <ul style="list-style-type: none"><li>• loose processor</li></ul>	
	wait(full) is now successful because something has been added to buffer


Silberschatz, Galvin and Gagne ©2005

## Readers-Writers

- Concurrent processes share a file, record, or other resources
- Some may read only (readers), some may both read and write (writers)
- Two concurrent reads have no adverse effects
- Problems if
  - concurrent reads and writes
  - multiple writes

Silberschatz, Galvin and Gagne ©2005


## Readers-Writers Diagram



Silberschatz, Galvin and Gagne ©2005

## The Readers-Writers Problem

Multiple readers or a single writer can use DB.



Silberschatz, Galvin and Gagne ©2005

## Readers-Writers Problem

- Two Variations
  - First Readers-Writers problem: No reader be kept waiting unless a writer has already obtained exclusive write permissions (Readers have high priority)
  - Second Readers-Writers problem: If a writer is waiting/ready, no new readers may start reading (Writers have high priority)
- A solution to either problem may result in starvation
  - In the first case, writers may starve; in the second case, reader may starve
- We present solution to first readers writers problem using semaphores

AE4B330SS

Lecture 6 / Page 49

Silberschatz, Galvin and Gagne ©2005

## Solution to First Readers-Writers Problem

- ❖ The reader processes share the following data structures:
 

```
semaphore mutex, wrt;
int readcount;
```
- ❖ The binary semaphores mutex and wrt are initialized to 1; readcount is initialized to 0;
- ❖ Semaphore wrt is common to both reader and writer process
  - ❖ wrt functions as a mutual exclusion for the writers
  - ❖ It is also used by the first or last reader that enters or exits the critical section
  - ❖ It is not used by readers who enter or exit while other readers are in their critical section
- ❖ The readcount variable keeps track of how many processes are currently reading the object
- ❖ The mutex semaphore is used to ensure mutual exclusion when readcount is updated

AE4B330SS

Lecture 6 / Page 50

Silberschatz, Galvin and Gagne ©2005

## Solution to first Readers-Writers Problem

- The structure of a writer process

```
do {
    wait (wrt) ;
    ...
    // writing is performed
    ...
    signal (wrt) ;
} while (TRUE);
```

AE4B330SS

Lecture 6 / Page 51

Silberschatz, Galvin and Gagne ©2005

## Solution to first Readers-Writers Problem

- The structure of a reader process

```
do {
    wait (mutex) ;
    readcount ++ ;
    if (readcount == 1)
        wait (wrt) ;
    signal (mutex)
    ...
    // reading is performed
    ...
    wait (mutex) ;
    readcount -- ;
    if (readcount == 0)
        signal (wrt) ;
    signal (mutex) ;
} while (TRUE);
```


AE4B330SS

Lecture 6 / Page 52

Silberschatz, Galvin and Gagne ©2005

## Dining Philosophers Problem

- Five philosophers sit at a round table - thinking and eating
- Each philosopher has one chopstick
  - five chopsticks total
- A philosopher needs two chopsticks to eat
  - philosophers must share chopsticks to eat
- No interaction occurs while thinking



AE4B33OSS

Lecture 6 / Page 53

Silberschatz, Galvin and Gagne ©2005

## Nature of Problem

- Problem:
  - starvation
    - a philosopher may never get the two chopsticks necessary to eat
  - deadlocks
    - two neighboring philosophers may try to eat at same time

4

AE4B33OSS

Lecture 6 / Page 54

Silberschatz, Galvin and Gagne ©2005

## Faulty Solution

Assume that no two neighbors are eating simultaneously

Shared data

binary semaphore chopstick [5];  
where all elements of chopstick are initialized to 1

```


■ The structure of Philosopher i:
do {
    wait(chopstick[i])
    wait(chopstick[(i+1) % 5])
    ...
    // eat
    ...
    signal(chopstick[i]);
    signal(chopstick[(i+1) % 5]);
    ...
    // think
    ...
} while (TRUE);
    
```

AE4B33OSS

Lecture 6 / Page 55

Silberschatz, Galvin and Gagne ©2005

## Deadlock



❖ Clearly deadlock is possible ( and starvation possible.)

❖ Several possible solutions for deadlock:

- ❖ Allow only 4 philosophers to be hungry at a time.
- ❖ Allow pickup only if both chopsticks are available. ( Done in critical section )
- ❖ Odd # philosopher always picks up left chopstick 1st, even # philosopher always picks up right chopstick 1st.

56

AE4B33OSS

Lecture 6 / Page 56

Silberschatz, Galvin and Gagne ©2005

### Problems with Semaphores

- Semaphores must be used with care. Incorrect use of semaphore results in following types of errors:
  - Inversed order: Suppose a process interchanges the order of wait() and signal() operations on the semaphore mutex
 

```

signal (mutex);
...
critical section
...
wait(mutex);
          
```

    - In this case, multiple processes may be executing in their critical sections, violating the mutual-exclusion requirement

AE4B330SS

Lecture 6 / Page 57

Silberschatz, Galvin and Gagne ©2005

### Problems with Semaphores

- Repeated calls: Suppose a process replaces signal(mutex) with wait(mutex)
 

```

wait(mutex);
...
critical section
...
wait(mutex);
      
```

  - In this case, processes may block forever (deadlock will occur)
- Omitted calls: Suppose a process omits the wait(mutex), or the signal (mutex) or both
  - In this case, either mutual exclusion is violated or a deadlock will occur

AE4B330SS

Lecture 6 / Page 58

Silberschatz, Galvin and Gagne ©2005

### Monitors

- Monitors provides a software solution to synchronization problems
- Monitor is a high-level abstraction that provides a convenient and effective mechanism for process synchronization

```

monitor monitor-name
{
  // shared variable declarations
  procedure P1 (...) { .... }
  ...
  procedure Pn (...) { ..... }
  initialization code ( .... ) { ... }
  ...
}
  
```


- Only one process may be active within the monitor at a time
- *Abstract data type*, internal variables only accessible by code within the procedure

AE4B330SS

Lecture 6 / Page 59

Silberschatz, Galvin and Gagne ©2005

### Schematic view of a Monitor



AE4B330SS

Lecture 6 / Page 60

Silberschatz, Galvin and Gagne ©2005

## Monitors


- To allow a process to wait within the monitor, one or more variables of type *condition* must be declared, as:  
condition x, y;
- Two operations on a condition variable:
  - **x.wait()** – a process invoking this operation is suspended until another process invokes **x.signal()**
  - **x.signal()** – resumes exactly one suspended process. If no process is suspended, then signal operation has no effect: the state of x is same as if the operation had never been executed

AE4B330SS

Lecture 6 / Page 61

Silberschatz, Galvin and Gagne ©2005

## Monitor with Condition Variables



AE4B330SS

Lecture 6 / Page 62

Silberschatz, Galvin and Gagne ©2005

## Condition Variables Choices

- If process P invokes **x.signal()**, with process Q in **x.wait()** state, what should happen next?
  - If Q is resumed, then P must wait. Otherwise, both P and Q would be active simultaneously within the monitor
- Two Options include
  - **Signal and wait** – P waits until Q leaves monitor or waits for another condition
  - **Signal and continue** – Q waits until P leaves the monitor or waits for another condition
- Both have pros and cons – language implementer can decide
- A compromise between these two choices was adopted in the Concurrent Pascal language
  - P executing signal immediately leaves the monitor, hence Q is resumed
- Implemented in other languages including Mesa, C#, Java

AE4B330SS

Lecture 6 / Page 63

Silberschatz, Galvin and Gagne ©2005

## Solution to Dining Philosophers using Monitors

- We need to distinguish among three states in which we may find a philosopher. For this purpose we introduce the following data structure:
 

```
enum{THINKING, HUNGRY, EATING} state[5]
```
- Philosopher i can set the variable `state[i] = EATING` only if her two neighbors are not eating:
 

```
(state[(i+4)%5] != EATING and (state[(i+1)%5] != EATING
```
- We also need to declare
 

```
condition self[5];
```

 in which philosopher i can delay herself when she is hungry but is unable to obtain the chopsticks she needs
- The distribution of the chopsticks is controlled by the monitor DiningPhilosophers, defined below:

AE4B330SS

Lecture 6 / Page 64

Silberschatz, Galvin and Gagne ©2005



### Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING, HUNGRY, EATING } state [5];
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

AE4B33OSS

Lecture 6 / Page 65

Silberschatz, Galvin and Gagne ©2005

### Solution to Dining Philosophers (Cont.)

```
void test (int i) {
    if ( (state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING;
        self[i].signal();
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
```

AE4B33OSS

Lecture 6 / Page 66

Silberschatz, Galvin and Gagne ©2005

### Solution to Dining Philosophers (Cont.)

- Each philosopher  $i$  invokes the operations `pickup()` and `putdown()` in the following sequence:

```
DiningPhilosophers.pickup(i);
...
eat
...
DiningPhilosophers.putdown(i);
```

- No deadlock, but starvation is possible

AE4B33OSS

Lecture 6 / Page 67

Silberschatz, Galvin and Gagne ©2005

### Monitor Implementation Using Semaphores

#### Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next_count = 0;
```

#### Each procedure $F$ will be replaced by

```
wait(mutex);
...
body of  $F$ 
...
if (next_count > 0)
    signal(next);
else
    signal(mutex);
```

- Mutual exclusion within a monitor is ensured.

AE4B33OSS

Lecture 6 / Page 68

Silberschatz, Galvin and Gagne ©2005

### Monitor Implementation Using Semaphores

- For each condition variable  $x$ , we have:  

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```
- The operation  $x.wait$  can be implemented as:  

```
x_count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x_count--;
```
- The operation  $x.signal$  can be implemented as:  

```
if (x_count > 0) {
    next_count++;
    signal(x_sem);
    wait(next);
    next_count--;
```

AE4B330SS

Lecture 6 / Page 69

Silberschatz, Galvin and Gagne ©2005

### Resuming Processes within a Monitor

- If several processes queued on condition  $x$ , and an  $x.signal()$  operation is executed by some process, then how do we determine which of the suspended processes which should be resumed next?
- FCFS frequently not adequate
- **conditional-wait** construct of the form  $x.wait(c)$ 
  - Where  $c$  is **priority number**
  - Process with lowest number (highest priority) is scheduled next
- ResourceAllocator monitor, which controls the allocation of a single resource among competing processes is given next.

AE4B330SS

Lecture 6 / Page 70

Silberschatz, Galvin and Gagne ©2005

### A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = FALSE;
    }
}
```

AE4B330SS

Lecture 6 / Page 71

Silberschatz, Galvin and Gagne ©2005

### A Monitor to Allocate Single Resource

- Each process, when requesting an allocation of this resource, specifies the maximum time it plans to use the resource
- Monitor allocates the resource to the process that has the shortest time-allocation request
- A process that needs to access the resource in question must observe the following sequence:  

```
R.acquire(t);
...
access the resource
...
R.release();
```

where  $R$  is an instance of type ResourceAllocator

AE4B330SS

Lecture 6 / Page 72

Silberschatz, Galvin and Gagne ©2005

### **Problems with Monitors**

- A process might access a resource without first gaining access permission to resource
- A process might never release a resource once it has been granted access to resource
- A process might attempt to release a resource that it never requested
- A process might request the same resource twice (without first releasing the resource)